

Decision Trees and Random Forests

Lecture 7 - DAMLF | ML1


Parametric vs Nonparametric Models

PARAMETRIC MODELS

Have a **fixed** number of parameters.

- + Simple
- + Fast
- Makes strong assumptions about data
- Prone to underfit*
**except high-capacity variants that depend on regularization*

Parametric vs Nonparametric Models




Linear Regression: $h(x; \theta) = \theta^T x$

PARAMETRIC MODELS

Have a **fixed** number of parameters.

- + Simple
- + Fast
- Makes strong assumptions about data
- Prone to underfit*
**except high-capacity variants that depend on regularization*

Parametric vs Nonparametric Models



Logistic Regression: $h(x; \theta) = g(\theta^T x)$

PARAMETRIC MODELS

Have a **fixed** number of parameters.

- + Simple
- + Fast
- Makes strong assumptions about data
- Prone to underfit*
**except high-capacity variants that depend on regularization*

Parametric vs Nonparametric Models

PARAMETRIC MODELS

Have a **fixed** number of parameters.

- + Simple
- + Fast
- Makes strong assumptions about data
- Prone to underfit*
**except high-capacity variants that depend on regularization*

NONPARAMETRIC MODELS

The number of parameters **grows** with the amount of training data.

- + Makes few if any assumptions about data
- + Can fit a larger class of functional forms
- Slow
- Prone to overfit

Parametric vs Nonparametric Models

PARAMETRIC MODELS

Have a **fixed** number of parameters.

- + Simple
- + Fast
- Makes strong assumptions about data
- Prone to underfit*
**except high-capacity variants that depend on regularization*

TWO NONPARAMETRIC MODELS

- **Support Vector Machines**
kernel method
- **Decision Trees**
adaptive basis function model

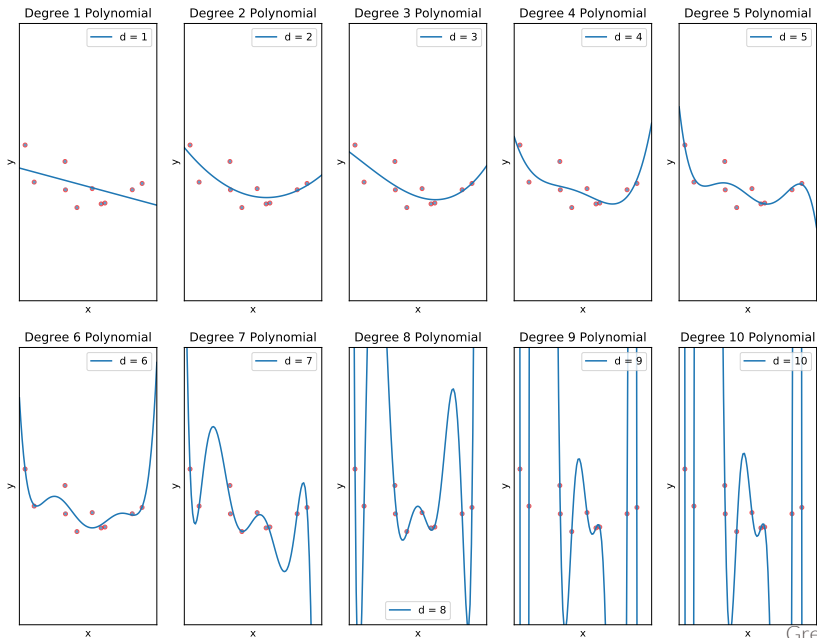
NONPARAMETRIC MODELS

The number of parameters **grows** with the amount of training data.

- + Makes few if any assumptions about data
- + Can fit a larger class of functional forms
- Slow
- Prone to overfit


Kernels and Adaptive Basis Functions

Review: Polynomial Regression



Review: Polynomial Regression

Degree 7 Polynomial



POLYNOMIAL REGRESSION

The degree 7 polynomial

$$h(x; \theta) = \theta_0 + \theta_1 x + \theta_2 x^2 + \theta_3 x^3 + \theta_4 x^4 + \theta_5 x^5 + \theta_6 x^6 + \theta_7 x^7$$


may be rewritten as

$$\begin{aligned} h(x; \theta) &= \theta^T \mathbf{z} \\ &= \theta_0 + \theta_1 \mathbf{z}_1 + \theta_2 \mathbf{z}_2 + \theta_3 \mathbf{z}_3 + \theta_4 \mathbf{z}_4 + \theta_5 \mathbf{z}_5 + \theta_6 \mathbf{z}_6 + \theta_7 \mathbf{z}_7 \end{aligned}$$

where $\mathbf{z} = [1, x, x^2, x^3, x^4, x^5, x^6, x^7]$.

Review: Polynomial Regression

Degree 7 Polynomial



POLYNOMIAL REGRESSION

The degree 7 polynomial

$$h(x; \theta) = \theta_0 + \theta_1 x + \theta_2 x^2 + \theta_3 x^3 + \theta_4 x^4 + \theta_5 x^5 + \theta_6 x^6 + \theta_7 x^7$$

may be rewritten as

$$\begin{aligned} h(x; \theta) &= \theta^T \mathbf{z} \\ &= \theta_0 + \theta_1 \mathbf{z}_1 + \theta_2 \mathbf{z}_2 + \theta_3 \mathbf{z}_3 + \theta_4 \mathbf{z}_4 + \theta_5 \mathbf{z}_5 + \theta_6 \mathbf{z}_6 + \theta_7 \mathbf{z}_7 \end{aligned}$$


where $\mathbf{z} = [1, x, x^2, x^3, x^4, x^5, x^6, x^7]$.

Suppose $f_j(x) = \mathbf{z}_j$ is a polynomial **basis function** of x whose range is $\mathbf{z}_j = x^j$. Then, we may rewrite $h(x; \theta)$ again as:

$$h(x; \theta) = \theta^T \mathbf{f}_j(x) \quad \text{for } j = 1, \dots, 7$$

Review: Polynomial Logistic Regression

NON-LINEAR DECISION BOUNDARIES



Review: Polynomial Logistic Regression


Problem: High Dimensionality.

Introducing non-linearity by explicitly engineering and computing a vector of new features for each input is impractical for multi-dimensional input.

Problem: Choosing f .

*In addition to the **curse of dimensionality**, you have to engineer the basis function, f .*

NON-LINEAR DECISION BOUNDARIES



KERNEL FUNCTION

A kernel method has the form

$$h(x) = \theta^T \phi(x)$$

where $\phi(x)$ is a list of **kernel functions** κ :

$$\phi(x) = [\kappa(x, \mu_1), \kappa(x, \mu_2), \dots, \kappa(x, \mu_m)]$$

Kernel Function

A MEASURE OF SIMILARITY

$$\phi(\mathbf{x}) = [\kappa(x, \mu_1), \kappa(x, \mu_2), \dots, \kappa(x, \mu_m)]$$

The **kernel function** $\kappa(x, \mu_k)$ is a real-valued function that is typically:

(non-negative): $\kappa(x, \mu_k) \geq 0$

(symmetric): $\kappa(x, \mu_k) = \kappa(\mu_k, x)$

Kernel Function

OK: High Dimensionality.

Problem: Choosing κ

A MEASURE OF SIMILARITY

$$\phi(\mathbf{x}) = [\kappa(\mathbf{x}, \mu_1), \kappa(\mathbf{x}, \mu_2), \dots, \kappa(\mathbf{x}, \mu_m)]$$

The **kernel function** $\kappa(\mathbf{x}, \mu_k)$ is a real-valued function that is typically:

(non-negative): $\kappa(\mathbf{x}, \mu_k) \geq 0$

(symmetric): $\kappa(\mathbf{x}, \mu_k) = \kappa(\mu_k, \mathbf{x})$

Kernel functions are interpreted as **measures of similarity** between training example \mathbf{x} and sample μ_k , where μ_k is either all m training examples or some (proper) subset of training examples.

Kernel Function

SOME KERNEL FUNCTIONS

- **Radial Basis (RBF)**

*Gaussian kernel
ARD kernel*

- **Cosine similarity**

to compare documents

- **Mercer kernels**

aka, positive definite kernels

- **Linear kernels**

for linearly separable features

- **Matern kernel**

Gaussian process regression

- **String kernels**

num. of strings in common

- **Pyramid match kernels**

*Mercer ker. for bag-of-words
representation of images*

- **Probability product**

probability generative model

- **Fisher kernels**

*cf. string kernels for Markov
chains*

A MEASURE OF SIMILARITY

$$\phi(\mathbf{x}) = [\kappa(\mathbf{x}, \mu_1), \kappa(\mathbf{x}, \mu_2), \dots, \kappa(\mathbf{x}, \mu_m)]$$

The **kernel function** $\kappa(\mathbf{x}, \mu_k)$ is a real-valued function that is typically:

(non-negative): $\kappa(\mathbf{x}, \mu_k) \geq 0$

(symmetric): $\kappa(\mathbf{x}, \mu_k) = \kappa(\mu_k, \mathbf{x})$

Kernel functions are interpreted as **measures of similarity** between training example \mathbf{x} and sample μ_k , where μ_k is either all m training examples or some (proper) subset of training examples.

Kernel Function

SOME KERNEL FUNCTIONS

- **Radial Basis (RBF)**

*Gaussian kernel
ARD kernel*

- **Cosine similarity**

to compare documents

- **Mercer kernels**

aka, positive definite kernels

- **Linear kernels**

for linearly separable features

- **Matern kernel**

Gaussian process regression

- **String kernels**

num. of strings in common

- **Pyramid match kernels**

*Mercer ker. for bag-of-words
representation of images*

- **Probability product**

probability generative model

- **Fisher kernels**

*cf. string kernels for Markov
chains*

TEMPLATE MATCHING

Kernel methods perform a type of **template matching**.

They compare the input x to **saved prototypes** μ_k .

Kernel Function

SOME KERNEL FUNCTIONS

- **Radial Basis (RBF)**

Gaussian kernel

ARD kernel

- **Cosine similarity**

to compare documents

- **Mercer kernels**

aka, positive definite kernels

- **Linear kernels**

for linearly separable features

- **Matern kernel**

Gaussian process regression

- **String kernels**

num. of strings in common

- **Pyramid match kernels**

*Mercer ker. for bag-of-words
representation of images*

- **Probability product**

probability generative model

- **Fisher kernels**

*cf. string kernels for Markov
chains*

TEMPLATE MATCHING

One way to **learn** the parameters of a kernel is to try the **ARD kernel**:

$$\kappa(\mathbf{x}, \mathbf{x}') = \theta_0 \exp \left(-\frac{1}{2} \sum_{j=1}^n \theta_j (x_j - x'_j)^2 \right)$$

where n is the total number of features, to estimate θ_j .

Kernel Methods

REQUIRE GOOD KERNELS

Kernel methods depend on having **good base kernels** and are **computationally expensive**.

REQUIRE GOOD KERNELS

Kernel methods depend on having **good base kernels** and are **computationally expensive**.

These limitations of kernel methods motivate other techniques that learn useful features from data.

Adaptive Basis Function Models

GET RID OF KERNELS

Instead, **learn** useful **features** $\phi(\mathbf{x})$ directly from the input data.

Adaptive Basis Function Models

GET RID OF KERNELS

Instead, **learn** useful **features** $\phi(\mathbf{x})$ directly from the input data.

Adaptive Basis Function models are of the form

$$h(\mathbf{x}) = \theta_0 + \sum_{k=1}^K \theta_k \phi_k(\mathbf{x})$$

where $\phi_k(\mathbf{x})$ is the k -th basis function learned from data.

Adaptive Basis Function Models (ABMs)

ABMs ARE NOT LINEAR IN PARAMETERS

$$h(\mathbf{x}) = \theta_0 + \sum_{k=1}^K \theta_k \phi_k(\mathbf{x})$$

Each basis function $\phi_k = \phi(\mathbf{x}, \boldsymbol{\omega}_k)$ is **parametric**.

Adaptive Basis Function Models (ABMs)

ABMs ARE NOT LINEAR IN PARAMETERS

$$h(\mathbf{x}) = \theta_0 + \sum_{k=1}^K \theta_k \phi_k(\mathbf{x})$$

Each basis function $\phi_k = \phi(\mathbf{x}, \boldsymbol{\omega}_k)$ is **parametric**.

Vector $\boldsymbol{\omega}_k$ parameterizes the basis function ϕ_k , and the full parameter set then is:

$$\boldsymbol{\theta} = [\theta_0, \theta_1, \dots, \theta_n, \{\boldsymbol{\omega}_1, \dots, \boldsymbol{\omega}_n\}]$$

Adaptive Basis Function Models (ABMs)

ABMs ARE NOT LINEAR IN PARAMETERS

$$h(\mathbf{x}) = \theta_0 + \sum_{k=1}^K \theta_k \phi_k(\mathbf{x})$$

Each basis function $\phi_k = \phi(\mathbf{x}, \boldsymbol{\omega}_k)$ is **parametric**.

Vector $\boldsymbol{\omega}_k$ parameterizes the basis function ϕ_k , and the full parameter set then is:

$$\boldsymbol{\theta} = [\theta_0, \theta_1, \dots, \theta_n, \{\boldsymbol{\omega}_1, \dots, \boldsymbol{\omega}_n\}]$$

But, the resulting model is **not linear** in $\boldsymbol{\theta}$.

Thus, ABMs are **nonparametric models**.

Adaptive Basis Function Models (ABMs)

SOME ADAPTIVE BASIS MODELS

- **Classification and Regression Trees (CART)**

*aka, decision trees;
random forests*

- **Generalized Additive Models (GAM)**

*multivariate adaptive
regression splines (MARS)*

- **Feed Forward Neural Networks**

aka, multi-layer perceptrons

- **Ensemble Learning**

ABMs ARE NOT LINEAR IN PARAMETERS

$$h(\mathbf{x}) = \theta_0 + \sum_{k=1}^K \theta_k \phi_k(\mathbf{x})$$

Each basis function $\phi_k = \phi(\mathbf{x}, \boldsymbol{\omega}_k)$ is **parametric**.

Vector $\boldsymbol{\omega}_k$ parameterizes the basis function ϕ_k , and the full parameter set then is:

$$\boldsymbol{\theta} = [\theta_0, \theta_1, \dots, \theta_n, \{\boldsymbol{\omega}_1, \dots, \boldsymbol{\omega}_n\}]$$

But, the resulting model is **not linear** in $\boldsymbol{\theta}$.

Thus, ABMs are **nonparametric models**.

Decision Trees

CART MODELS

Classification and Regression Trees (CART) are adaptive basis function models.

CART models **recursively partition** the input, with each cell of the partition covering a region of the input.

Then, a **local model** is constructed for each of the resulting regions.

Decision Trees

CART MODELS


A CART model can be represented by a **tree**, where leaves correspond to the partitioned regions of the input space.

Decision Trees


CART MODELS

A CART model can be represented by a **tree**, where leaves correspond to the partitioned regions of the input space.

Example: two inputs, x_1 and x_2 .



Decision Trees




Based on Murphy (2012, Figure 16.1)


CART MODELS

A CART model can be represented by a **tree**, where leaves correspond to the partitioned regions of the input space.

Example: two inputs, x_1 and x_2 .



Decision Trees




Based on Murphy (2012, Figure 16.1)

CART MODELS

In this example, the **2d space** is partitioned into 5 regions: R_1 , R_2 , R_3 , R_4 , R_5 .

Partitions splits are **parallel to the axes** x_1, x_2 .

Decision Trees



Based on Murphy (2012, Figure 16.1)

MEAN RESPONSE IN EACH REGION

Define $\mathbf{1}_R(x) = 1$ if $x \in R$ and 0 otherwise.

Then, the equation


$$h(\mathbf{x}) = \mathbb{E}[y | \mathbf{x}] = \sum_{k=1}^K \theta_k \mathbf{1}_{R_k}(\mathbf{x}) = \sum_{k=1}^K \theta_k \phi(\mathbf{x}, \omega_k)$$

specifies a **CART model**.

Where:

- R_k is the k 'th region,
- θ_k is the **mean response** in R_k
- ω_k encodes both the **choice of variables to split** and the **threshold values** for splitting each variable along the path from the root to the k 'th leaf.

Decision Trees



Based on Murphy (2012, Figure 16.1)

MEAN RESPONSE IN EACH REGION

Define $\mathbf{1}_R(\mathbf{x}) = 1$ if $\mathbf{x} \in R$ and 0 otherwise.

Then, the equation

$$h(\mathbf{x}) = \mathbb{E}[y | \mathbf{x}] = \sum_{k=1}^K \theta_k \mathbf{1}_{R_k}(\mathbf{x}) = \sum_{k=1}^K \theta_k \phi(\mathbf{x}, \omega_k)$$

specifies a **CART model**.

CART is an **adaptive basis-function model**, where

- **basis functions** define regions;
- **parameters** specify the response value in each region.

CART Classification

CLASSIFICATION

Instead of storing the mean response at each leaf, store the **distribution over class labels** at each leaf.


CART Classification

CLASSIFICATION

Instead of storing the mean response at each leaf, store the **distribution over class labels** at each leaf.

Example:

- Suppose x matches the color criterion **blue**.
What is the probability that $y = 1$?



Example from Murphy (2012, Figure 1.1)

CART Classification


CLASSIFICATION

Instead of storing the mean response at each leaf, store the **distribution over class labels** at each leaf.

Example:


- Suppose x matches the color criterion **blue**.
What is the probability that $y = 1$?

$$p(y = 1 \mid x \in \text{blue}) = 1$$



Example from Murphy (2012, Figure 1.1)

CART Classification



Example from Murphy (2012, Figure 1.1)

CLASSIFICATION

Instead of storing the mean response at each leaf, store the **distribution over class labels** at each leaf.


Example:

- Suppose x matches the color criterion **blue**. What is the probability that $y = 1$?

$$p(y = 1 \mid x \in \text{blue}) = 1$$

- Suppose x matches the color criterion **red** and the shape criterion **ellipse**. What is the probability that $y = 1$?

CART Classification



Example from Murphy (2012, Figure 1.1)

CLASSIFICATION

Instead of storing the mean response at each leaf, store the **distribution over class labels** at each leaf.

Example:

- Suppose x matches the color criterion **blue**. What is the probability that $y = 1$?

$$p(y = 1 \mid x \in \text{blue}) = 1$$


- Suppose x matches the color criterion **red** and the shape criterion **ellipse**. What is the probability that $y = 1$?

$$p(y = 1 \mid x \in \text{red} \cap \text{ellipse}) = \frac{1}{2}$$


CART Classification

CLASSIFICATION

Instead of storing the mean response at each leaf, store the **distribution over class labels** at each leaf.




Example from Murphy (2012, Figure 1.1)




CART Classification

CLASSIFICATION

Instead of storing the mean response at each leaf, store the **distribution over class labels** at each leaf.




Example from Murphy (2012, Figure 1.1)



Each leaf (n_1, n_0) says there are n_1 **positive** examples in that path that match and n_0 **negative** examples that do not match.

Hypothesis Representation




CART:

Hypothesis h computes the **expectation** of y given features \mathbf{x} :

$$h(\mathbf{x}) = \mathbb{E}[y | \mathbf{x}] = \sum_{k=1}^K \theta_k \mathbf{1}_{R_k}(\mathbf{x}) = \sum_{k=1}^K \theta_k \phi(\mathbf{x}, \omega_k)$$

Hypothesis Representation




CART:

Hypothesis h computes the **expectation** of y given features \mathbf{x} :

$$h(\mathbf{x}) = \mathbb{E}[y | \mathbf{x}] = \sum_{k=1}^K \theta_k \mathbf{1}_{R_k}(\mathbf{x}) = \sum_{k=1}^K \theta_k \phi(\mathbf{x}, \omega_k)$$

The form of the hypothesis depends on what variables are **split** at which **threshold**, which is encoded in ω_k for each $k \in K$ leaf.

Hypothesis Representation



CART:

Hypothesis h computes the **expectation** of y given features \mathbf{x} :

$$h(\mathbf{x}) = \mathbb{E}[y | \mathbf{x}] = \sum_{k=1}^K \theta_k \mathbf{1}_{R_k}(\mathbf{x}) = \sum_{k=1}^K \theta_k \phi(\mathbf{x}, \omega_k)$$

The form of the hypothesis depends on what variables are **split** at which **threshold**, which is encoded in ω_k for each $k \in K$ leaf.

Each ω_k is selected by a **split function**.

Growing a Tree

OPTIMAL PARTITIONS

The **split function** chooses the best feature j^* and the best threshold value t^* for that feature, as follows:

$$(j^*, t^*) = \arg \min_{\substack{j \in \{1, \dots, n\} \\ t \in \mathcal{T}_j}} \min \text{Cost}(\{\mathbf{x}_i, y_i : x_{i,j} \leq t\}) + \text{Cost}(\{\mathbf{x}_i, y_i : x_{i,j} > t\})$$

where:

- n is the dimension of the features
- \mathcal{T}_j is the set of possible thresholds for feature j
- **Cost** is a cost function, to be defined
- j^* is the best feature from the n features
- t^* is the best threshold value for feature j^*

Cost Functions

$$(j^*, t^*) = \arg \min_{j \in \{1, \dots, n\}} \min_{t \in \mathcal{T}_j} \text{Cost}(\{\mathbf{x}_i, y_i : x_{i,j} \leq t\}) + \text{Cost}(\{\mathbf{x}_i, y_i : x_{i,j} > t\})$$

Regression Cost Function:

$$\text{Cost}(\mathcal{D}) = \sum_{i=1}^{|\mathcal{D}|} (\bar{y} - y_i)^2$$

where \mathcal{D} is the specified data set, $|\mathcal{D}| = m$, and \bar{y} is the **mean** of the target y within the set of data \mathcal{D} :

$$\bar{y} = \frac{1}{m} \sum_{i=1}^m y_i$$

Cost Functions

$$(j^*, t^*) = \arg \min_{j \in \{1, \dots, n\}} \min_{t \in \mathcal{T}_j} \text{Cost}(\{\mathbf{x}_i, y_i : x_{i,j} \leq t\}) + \text{Cost}(\{\mathbf{x}_i, y_i : x_{i,j} > t\})$$

Regression Cost Function:

$$\text{Cost}(\mathcal{D}) = \sum_{i=1}^{|\mathcal{D}|} (\bar{y} - y_i)^2$$

where \mathcal{D} is the specified data set, $|\mathcal{D}| = m$, and \bar{y} is the **mean** of the target y within the set of data \mathcal{D} :

$$\bar{y} = \frac{1}{m} \sum_{i=1}^m y_i$$

Alternatively, a linear regression model can be **fit to each leaf** using the features selected for that path as input, then residual error can be measured.

Cost Functions

$$(j^*, t^*) = \arg \min_{\substack{j \in \{1, \dots, n\} \\ t \in \mathcal{T}_j}} \min \text{Cost}(\{\mathbf{x}_i, y_i : x_{i,j} \leq t\}) + \text{Cost}(\{\mathbf{x}_i, y_i : x_{i,j} > t\})$$

Classification:

Step 1.

Compute the proportion of class c observations in leaf j passing test $X_j < t$, written $\hat{\pi}_{c,j}$, by

$$\hat{\pi}_{c,j} = \frac{1}{|\mathcal{D}_j|} \sum_{i=1}^{|\mathcal{D}_j|} \mathbf{1}(y_i = c)$$

where \mathcal{D}_j is the data in leaf j and the indicator $\mathbf{1}(y_i = c) = 1$, if y_i has class label c , 0 otherwise.

Cost Functions

$$(j^*, t^*) = \arg \min_{\substack{j \in \{1, \dots, n\} \\ t \in \mathcal{T}_j}} \min \text{Cost}(\{\mathbf{x}_i, y_i : x_{i,j} \leq t\}) + \text{Cost}(\{\mathbf{x}_i, y_i : x_{i,j} > t\})$$

Classification:

Step 1.

Compute the **proportion** of class c observations in leaf j passing test $X_j < t$, written $\hat{\pi}_{c,j}$, by

$$\hat{\pi}_{c,j} = \frac{1}{|\mathcal{D}_j|} \sum_{i=1}^{|\mathcal{D}_j|} \mathbf{1}(y_i = c)$$

where \mathcal{D}_j is the data in leaf j and the **indicator** $\mathbf{1}(y_i = c) = 1$, if y_i has class label c , 0 otherwise.

$\hat{\pi}_{c,j}$ is a **Maximum Likelihood Estimator** for the distribution $p(c \mid X_j < t)$

Cost Functions

$$(j^*, t^*) = \arg \min_{\substack{j \in \{1, \dots, n\} \\ t \in \mathcal{T}_j}} \min \text{Cost}(\{\mathbf{x}_i, y_i : x_{i,j} \leq t\}) + \text{Cost}(\{\mathbf{x}_i, y_i : x_{i,j} > t\})$$

Classification:

Step 2.

There are many **cost functions** to measure the **quality of a split** in a CART classification model.

Cost Functions

$$(j^*, t^*) = \arg \min_{\substack{j \in \{1, \dots, n\} \\ t \in \mathcal{T}_j}} \min \text{Cost}(\{\mathbf{x}_i, y_i : x_{i,j} \leq t\}) + \text{Cost}(\{\mathbf{x}_i, y_i : x_{i,j} > t\})$$

Classification:

Step 2.

There are many **cost functions** to measure the **quality of a split** in a CART classification model.

Three common error measures for evaluating a proposed partition:

1. **Misclassification rate**
2. **Entropy**
3. **Gini index**

Misclassification Rate

1. MISCLASSIFICATION ERROR

Let $\hat{y}_c = \arg \max_c \hat{\pi}_c$ denote the **most probable class** label c for some (index omitted) region.

The **misclassification rate** error measure is then

$$\frac{1}{|\mathcal{D}_j|} \sum_{i=1}^{|\mathcal{D}_j|} \mathbf{1}(y_i \neq \hat{y}_{c,j}) = 1 - \hat{\pi}_{\hat{y},j}$$

Entropy

2. ENTROPY

The **entropy** measure is

$$-\sum_{c=1}^C \hat{\pi}_c \log \hat{\pi}_c$$

Entropy

2. ENTROPY

The **entropy** measure is

$$\mathbb{H}(\hat{\pi}) = - \sum_{c=1}^C \hat{\pi}_c \log \hat{\pi}_c$$

Entropy

2. ENTROPY

The **entropy** measure is

$$\mathbb{H}(\hat{\pi}) = - \sum_{c=1}^C \hat{\pi}_c \log \hat{\pi}_c$$

Minimizing entropy is equivalent to **maximizing information gain** between the test $X_j < t$ and the class label $Y = y$

Entropy

2. ENTROPY

The **entropy** measure is

$$\mathbb{H}(\hat{\pi}) = - \sum_{c=1}^C \hat{\pi}_c \log \hat{\pi}_c$$

Minimizing entropy is equivalent to **maximizing information gain** between the test $X_j < t$ and the class label $Y = y$

INFORMATION GAIN

The **info gain** of test $X_j < t$ wrt to label Y :

$$\text{InfoGain}(X_j < t, Y) := \mathbb{H}(Y) - \mathbb{H}(Y|X_j < t)$$

Entropy

2. ENTROPY

The **entropy** measure is

$$\mathbb{H}(\hat{\pi}) = - \sum_{c=1}^C \hat{\pi}_c \log \hat{\pi}_c$$

Minimizing entropy is equivalent to **maximizing information gain** between the test $X_j < t$ and the class label $Y = y$

INFORMATION GAIN

The **info gain** of test $X_j < t$ wrt to label Y :

$$\text{InfoGain}(X_j < t, Y) := \mathbb{H}(Y) - \mathbb{H}(Y|X_j < t)$$

where:

$$\begin{aligned} \text{InfoGain}(X_j < t, Y) = & \left(- \sum_c p(y = c) \log p(y = c) \right) \\ & + \left(- \sum_c p(y = c|X_j < t) \log p(y = c|X_j < t) \right) \end{aligned}$$

since $\hat{\pi}$ is an **maximum likelihood estimate** for the conditional distribution

$$p(c|X_j < t).$$

3. GINI INDEX


The **Gini index**

$$\sum_{c=1}^C \hat{\pi}_c (1 - \hat{\pi}_c) = 1 - \sum_{c=1}^C \hat{\pi}_c^2$$

is the **expected error rate**, where

- $\hat{\pi}_c$ is the probability a random entry in a leaf belongs to class c
- $(1 - \hat{\pi}_c)$ is the probability of being misclassified to class c

Comparison



TWO-CLASS EXAMPLE

Suppose $c = 0$ or $c = 1$ and $p = \hat{\pi}_{c=1}$. Then,

Misclassification Rate: $1 - \max(p, 1 - p)$

Entropy: $-p \log p - (1 - p) \log(1 - p)$


- range is 0 to 1.

Gini index: $p(1 - p)$

- range is 0 to $1/2$.

¹Entropy is scaled in the graph to pass through 0.5

Comparison



TWO-CLASS EXAMPLE

Suppose $c = 0$ or $c = 1$ and $p = \hat{\pi}_{c=1}$. Then,

Misclassification Rate: $1 - \max(p, 1 - p)$

Entropy: $-p \log p - (1 - p) \log(1 - p)$

- range is 0 to 1.

Gini index: $p(1 - p)$

- range is 0 to $1/2$.

Gini and Entropy¹ are very similar. Both are more sensitive to changes in class probability than misclassification.


¹Entropy is scaled in the graph to pass through 0.5

Comparison


EXAMPLE WITH FOUR OBSERVATIONS

Entropy: $-p \log p - (1 - p) \log(1 - p)$

Gini index: $p(1 - p)$



Comparison



EXAMPLE WITH FOUR OBSERVATIONS

Entropy: $-p \log p - (1-p) \log(1-p)$

Gini index: $p(1-p)$

PPPP: $p = 4/4$

Entropy: $-\frac{4}{4} \log \frac{4}{4} - (1 - \frac{4}{4}) \log(1 - \frac{4}{4}) = 0$


Gini: $\frac{4}{4}(1 - \frac{4}{4}) = 0$

PPNN: $p = 2/4$

Entropy: $-\frac{1}{2} \log \frac{1}{2} - (1 - \frac{1}{2}) \log(1 - \frac{1}{2}) = 1.0$

Gini: $2 \cdot \frac{1}{2}(1 - \frac{1}{2}) = 1.0$ # rescaled to $[0,1]$

Comparison



EXAMPLE WITH FOUR OBSERVATIONS

Entropy: $-p \log p - (1-p) \log(1-p)$

Gini index: $p(1-p)$

PPPP: $p = 4/4$

Entropy: $-\frac{4}{4} \log \frac{4}{4} - (1 - \frac{4}{4}) \log(1 - \frac{4}{4}) = 0$

Gini: $\frac{4}{4}(1 - \frac{4}{4}) = 0$

PPPN: $p = 3/4$

Entropy: $-\frac{3}{4} \log \frac{3}{4} - (1 - \frac{3}{4}) \log(1 - \frac{3}{4}) \approx 0.81$


Gini: $2 \cdot \frac{3}{4}(1 - \frac{3}{4}) = 0.75$ # rescaled to $[0,1]$

PPNN: $p = 2/4$

Entropy: $-\frac{1}{2} \log \frac{1}{2} - (1 - \frac{1}{2}) \log(1 - \frac{1}{2}) = 1.0$


Gini: $2 \cdot \frac{1}{2}(1 - \frac{1}{2}) = 1.0$ # rescaled to $[0,1]$

Overfitting



CART MODELS TEND TO OVERFIT


Overfitting



CART MODELS TEND TO OVERFIT

This CART classification model has **low** training error but **high** CV error.


Overfitting



CART MODELS TEND TO OVERFIT

This CART classification model has **low** training error but **high** CV error.

Overfitting



CART MODELS TEND TO OVERFIT

By reducing the **maximum tree depth**, you can address overfitting.

For example, adjusting the hyperparameter **max_depth** changes the regions as follows.

» max_depth = 5


max_depth = 4

max_depth = 3

max_depth = 2

max_depth = 1

Overfitting



CART MODELS TEND TO OVERFIT

By reducing the **maximum tree depth**, you can address overfitting.

For example, adjusting the hyperparameter **max_depth** changes the regions as follows.

max_depth = 5


» max_depth = 4

max_depth = 3

max_depth = 2

max_depth = 1

Overfitting



CART MODELS TEND TO OVERFIT

By reducing the **maximum tree depth**, you can address overfitting.

For example, adjusting the hyperparameter **max_depth** changes the regions as follows.

`max_depth = 5`


`max_depth = 4`

» `max_depth = 3`

`max_depth = 2`

`max_depth = 1`

Overfitting



CART MODELS TEND TO OVERFIT

By reducing the **maximum tree depth**, you can address overfitting.

For example, adjusting the hyperparameter **max_depth** changes the regions as follows.

max_depth = 5


max_depth = 4

max_depth = 3

» max_depth = 2

max_depth = 1

Overfitting



CART MODELS TEND TO OVERFIT

By reducing the **maximum tree depth**, you can address overfitting.

For example, adjusting the hyperparameter **max_depth** changes the regions as follows.

max_depth = 5


max_depth = 4

max_depth = 3

max_depth = 2

>> max_depth = 1

Overfitting



max_depth = 2

CART MODELS TEND TO OVERFIT

The same applies to CART regression models.


» max_depth = 2

max_depth = 3

max_depth = 5

max_depth = 10

Overfitting



$\text{max_depth} = 3$

CART MODELS TEND TO OVERFIT

The same applies to CART regression models.


$\text{max_depth} = 2$

» $\text{max_depth} = 3$

$\text{max_depth} = 5$

$\text{max_depth} = 10$

Overfitting



max_depth = 5

CART MODELS TEND TO OVERFIT

The same applies to CART regression models.


max_depth = 2

max_depth = 3

>> max_depth = 5

max_depth = 10

Overfitting



max_depth = 10

CART MODELS TEND TO OVERFIT

The same applies to CART regression models.


max_depth = 2

max_depth = 3

max_depth = 5

>> max_depth = 10

Overfitting



max_depth = 10

CART MODELS TEND TO OVERFIT

The same applies to CART regression models.

max_depth = 2

max_depth = 3

max_depth = 5

>> max_depth = 10

Technical note: The standard approach is to grow a full tree, then **prune** the tree back to the maximum depth.

Advantages and Disadvantages

CART Pros

- + Easy to interpret
- + Easily handle discrete and continuous input
- + Insensitive to monotone transformations of input
- + Perform automatic variable selection
- + Scale reasonably well
- + Relatively robust to outliers
- + Can be modified to handle missing inputs by looking for “backup” variables.

CART Cons

- Poor prediction accuracy
- Instability: small changes to input data can have large effects on structure of tree

Random Forests

BOOTSTRAP AGGREGATING

Trees are **high variance** estimators.

One technique for reducing variance of an estimate is to **average together** several estimates


BOOTSTRAP AGGREGATING

Suppose we train K different trees on subsets of data, randomly drawn with replacement, then calculate the **ensemble** h :

$$h(\mathbf{x}) = \sum_{t=1}^K \frac{1}{K} h_k(\mathbf{x})$$

where $h_k(\mathbf{x})$ is the k 'th tree.

Bagging



BOOTSTRAP AGGREGATING

Suppose we train K different trees on subsets of data, randomly drawn with replacement, then calculate the **ensemble** h :


$$h(\mathbf{x}) = \sum_{t=1}^K \frac{1}{K} h_k(\mathbf{x})$$

where $h_k(\mathbf{x})$ is the k 'th tree.

The result $h(\mathbf{x})$:

- combines the results (regression), or
- tallies votes for each class (classification).

Bagging



$h(\mathbf{x}) = \text{BLUE}$

BOOTSTRAP AGGREGATING

Suppose we train K different trees on subsets of data, randomly drawn with replacement, then calculate the **ensemble** h :

$$h(\mathbf{x}) = \sum_{t=1}^K \frac{1}{K} h_k(\mathbf{x})$$

where $h_k(\mathbf{x})$ is the k 'th tree.

The result $h(\mathbf{x})$:

- combines the results (regression), or
- tallies votes for each class (classification).

This method is called **bagging**, which is short for 'bootstrap aggregating'.

BOOTSTRAP AGGREGATING FOR CLASSIFICATION

Bagging creates a **committee** of trees to vote on a class, and the predicted class is the majority winner.

The idea behind bagging is to average many noisy but approximately unbiased models.

BOOTSTRAP AGGREGATING FOR CLASSIFICATION

Bagging creates a **committee** of trees to vote on a class, and the predicted class is the majority winner.

The idea behind bagging is to average many noisy but approximately unbiased models.

However, one **problem** running the same algorithm on different samples of data is **highly correlated predictors**, which limits the amount of variance reduction.

Random Forests

DE-CORRELATED TREES

Random forests:

- randomly sample a **subset of data** (like bagging), and
- randomly sample a **subset of input variables**.

random forests build a large collection of de-correlated trees and averages them.

Random Forests

Random Forests for Regression or Classification

- 1: **for** $b = 1$ to B **do**:
- 2: Draw a bootstrap sample Z of size N from training data
- 3: Grow a random forest tree T_b to the bootstrapped data:
- 4: **repeat**
- 5: (i) Select k variables at random from the n variables
- 6: (ii) Pick the best variable/split-thresholds among the k
- 7: (iii) Split the node into two daughter nodes.
- 8: **until** the minimum node size n_{\min} is reached.
- 9: **end for**
- 10: Output the ensemble of trees $\{T_1, T_2, \dots, T_B\}$.

To make a prediction given a new datapoint x :

Regression: $\frac{1}{B} \sum_{b=1}^B T_b(x)$

Classification: Let $\hat{\pi}_c(x)$ be the class prediction of the b 'th random-forest tree. Then the random forest prediction, $\hat{\pi}_{rf}(x)$, is the **majority vote** of

$$\{\hat{\pi}_1(x), \hat{\pi}_2(x), \dots, \hat{\pi}_b(x), \dots, \hat{\pi}_B(x)\}$$

References

Hastie, T., R. Tibshirani, and J. Friedman (2009).
The Elements of Statistical Learning (2nd ed.).
Springer Series in Statistics. Springer.

Murphy, K. P. (2012).
Machine Learning: A Probabilistic Perspective.
Cambridge, MA: MIT Press.