


# Neural Network Cost Function and Backpropagation

Lecture 9 - DAMLF | ML1

# Review: Forward Propagation



Equations to plug in:

$$z^{(2)} = \begin{bmatrix} z_1^{(2)} \\ z_2^{(2)} \\ z_3^{(2)} \end{bmatrix} \quad a^{(1)} = \begin{bmatrix} a_0^{(1)} \\ a_1^{(1)} \\ a_2^{(1)} \\ a_3^{(1)} \end{bmatrix}$$

$$a_1^{(2)} = g \left( \Theta_{10}^{(1)} a_0^{(1)} + \Theta_{11}^{(1)} a_1^{(1)} + \Theta_{12}^{(1)} a_2^{(1)} + \Theta_{13}^{(1)} a_3^{(1)} \right)$$

$$a_2^{(2)} = g \left( \Theta_{20}^{(1)} a_0^{(1)} + \Theta_{21}^{(1)} a_1^{(1)} + \Theta_{22}^{(1)} a_2^{(1)} + \Theta_{23}^{(1)} a_3^{(1)} \right)$$

$$a_3^{(2)} = g \left( \Theta_{30}^{(1)} a_0^{(1)} + \Theta_{31}^{(1)} a_1^{(1)} + \Theta_{32}^{(1)} a_2^{(1)} + \Theta_{33}^{(1)} a_3^{(1)} \right)$$

$$h(x^{(l)}; \theta) = g \left( \Theta_{10}^{(2)} a_0^{(2)} + \Theta_{11}^{(2)} a_1^{(2)} + \Theta_{12}^{(2)} a_2^{(2)} + \Theta_{13}^{(2)} a_3^{(2)} \right)$$

$$z^{(2)} = \Theta^{(1)} a^{(1)}$$



$$a^{(2)} = g \left( z^{(2)} \right)$$

$$z^{(3)} = \Theta^{(2)} a^{(2)}$$

$$a^{(3)} = h(x^{(l)}; \theta) = g \left( z^{(3)} \right)$$

$$a_0^{(2)} = 1 = a_0^{(1)} \text{ \# bias units}$$

# Review: Neural Network Architecture



# Multi-class Classification with Neural Networks

# One-vs-all Networks with Multiple Outputs



fish



amphibians



birds



reptiles



mammals

# One-vs-all Networks with Multiple Outputs



fish



amphibians




birds



reptiles



mammals



$$h(x^{(i)}; \theta) \in \mathbb{R}^5$$

# One-vs-all Networks with Multiple Outputs



fish

$$h(x^{(i)}; \theta) = \begin{bmatrix} 1 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$



amphibians

$$h(x^{(i)}; \theta) = \begin{bmatrix} 0 \\ 1 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$



birds

$$h(x^{(i)}; \theta) = \begin{bmatrix} 0 \\ 0 \\ 1 \\ 0 \\ 0 \end{bmatrix}$$



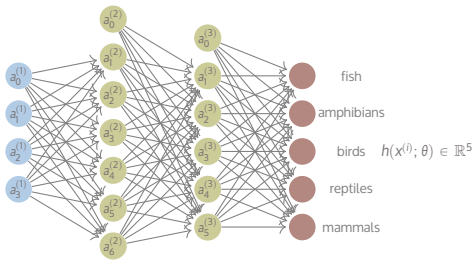
reptiles


$$h(x^{(i)}; \theta) = \begin{bmatrix} 0 \\ 0 \\ 0 \\ 1 \\ 0 \end{bmatrix}$$



mammals

$$h(x^{(i)}; \theta) = \begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 1 \end{bmatrix}$$






**Training Set:**  $(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})$ , where each  $y^{(i)}$  is **one-hot encoded** as

$$y^{(i)} = \begin{bmatrix} 1 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix} \text{ or } \begin{bmatrix} 0 \\ 1 \\ 0 \\ 0 \\ 0 \end{bmatrix} \text{ or } \begin{bmatrix} 0 \\ 0 \\ 1 \\ 0 \\ 0 \end{bmatrix} \text{ or } \begin{bmatrix} 0 \\ 0 \\ 0 \\ 1 \\ 0 \end{bmatrix} \text{ or } \begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 1 \end{bmatrix}$$

**Goal:**  $h(x^{(i)}; \theta) \approx y^{(i)}$



# Multi-class Classification Notation




**Training set:**  $\{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})\}$

**Notation:**

$m$	=	Number of training examples
$L$	=	Number of layers in the network

# Multi-class Classification Notation




**Training set:**  $\{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})\}$

**Notation:**

$m$	=	Number of training examples
$L$	=	Number of layers in the network
$a^{(1)}$	=	Input variables / Input layer
$a^{(2)}, \dots, a^{(L-1)}$	=	Hidden layers
$a^{(L)}$	=	Output layer
$s_l$	=	Number of units ( <b>not</b> counting the bias unit) in layer $l$

# Multi-class Classification Notation



**Training set:**  $\{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})\}$

**Example:**

$m$  = Number of training examples

$L$  = 4

$s_1$  = 3


$s_2$  = 6

$s_3$  = 5

$s_4$  = 5

# Binary & Multi-class Classification


## Binary Classification



1 output unit:  $h(x^{(i)}; \theta) \in \mathbb{R}$

$y = 1$  or  $y = 0$

## Multi-class Classification (K-classes)




$K$  output units:  $h(x^{(i)}; \theta) \in \mathbb{R}^K$

$$y^{(i)} = \begin{bmatrix} 1 \\ 0 \\ 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 1 \\ 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 0 \\ 1 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 0 \\ 0 \\ 1 \end{bmatrix}, \text{ for } K = 4.$$

# Binary & Multi-class Classification

## Binary Classification




1 output unit:  $h(x^{(i)}; \theta) \in \mathbb{R}$

$y = 1$  or  $y = 0$

$s_L = 1$  (and  $K = 1$ )

## Multi-class Classification (K-classes)



$K$  output units:  $h(x^{(i)}; \theta) \in \mathbb{R}^K$

$y^{(i)} = \begin{bmatrix} 1 \\ 0 \\ 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 1 \\ 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 0 \\ 1 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 0 \\ 0 \\ 1 \end{bmatrix}$ , for  $K = 4$ .

$s_L = K$  (and  $K > 2$ )

## Cost Function for Logistic Regression

# Review: Logistic Regression Cost Function

## Regularized Cost Function:

$$J(\Theta) = -\frac{1}{m} \left( \sum_{i=1}^m y^{(i)} \log h(x^{(i)}; \theta) + (1 - y^{(i)}) \log (1 - h(x^{(i)}; \theta)) \right) + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2$$

# Cost Function for Neural Networks

**Logistic Regression Cost Function for  $h(x^{(i)}; \Theta) \in \mathbb{R}$ :**

$$J(\Theta) = -\frac{1}{m} \left( \sum_{i=1}^m y^{(i)} \log h(x^{(i)}; \theta) + (1 - y^{(i)}) \log (1 - h(x^{(i)}; \theta)) \right) + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2$$

$$J(\Theta) = -\frac{1}{m} ( \text{Cost Function} ) + \text{Regularization Term}$$



# Cost Function for Neural Networks

**Logistic Regression Cost Function for  $h(x^{(i)}; \theta) \in \mathbb{R}$ :**

$$J(\theta) = -\frac{1}{m} \left( \sum_{i=1}^m y^{(i)} \log h(x^{(i)}; \theta) + (1 - y^{(i)}) \log (1 - h(x^{(i)}; \theta)) \right) + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2$$

$$J(\theta) = -\frac{1}{m} ( \text{Cost Function} ) + \text{Regularization Term}$$

**Neural Network Cost Function for  $h(x^{(i)}; \theta) \in \mathbb{R}^K$**

$$J(\theta) = -\frac{1}{m} ( \text{K Cost Functions} ) + \text{L-1 layers of Regularization Terms}$$

# Cost Function for Neural Networks

**Neural Network Cost Function for**  $h(x^{(i)}; \theta) \in \mathbb{R}^K$

$$J(\theta) = -\frac{1}{m} (\text{K Cost Functions}) + \text{L-1 layers of Regularization Terms}$$

# Cost Function for Neural Networks

**Neural Network Cost Function for**  $h(x^{(i)}; \Theta) \in \mathbb{R}^K$

$$J(\Theta) = -\frac{1}{m} (\text{K Cost Functions}) + \text{L-1 layers of Regularization Terms}$$

$$J(\Theta) = -\frac{1}{m} \left( \sum_{i=1}^m \sum_{k=1}^K y_k^{(i)} \log(h(x^{(i)}; \Theta)_k) + (1 - y_k^{(i)}) \log(1 - h(x^{(i)}; \Theta)_k) \right) + \frac{\lambda}{2m} \sum_{l=1}^{L-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} (\Theta_{ij}^{(l)})^2$$

# Cost Function for Neural Networks

**Neural Network Cost Function for**  $h(x^{(i)}; \Theta) \in \mathbb{R}^K$

$$J(\Theta) = -\frac{1}{m} (\text{K Cost Functions}) + \text{L-1 layers of Regularization Terms}$$

$$J(\Theta) = -\frac{1}{m} \left( \sum_{i=1}^m \sum_{k=1}^K y_k^{(i)} \log(h(x^{(i)}; \Theta)_k) + (1 - y_k^{(i)}) \log(1 - h(x^{(i)}; \Theta)_k) \right) + \frac{\lambda}{2m} \sum_{l=1}^{L-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} (\Theta_{ij}^{(l)})^2$$

where  $h(x^{(i)}; \Theta)_k$  is the  $k$ th output node of the  $\mathbb{R}^K$  output vector, and

$$\sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} (\Theta_{ij}^{(l)})^2 = \|\Theta^{(l)}\|_F^2$$

is the **Frobenius norm**, i.e., *the sum of squared elements of a matrix*.

Note that  $\Theta_{i,0}$  is not included in the regularization terms.

# Backpropagation for Minimizing the Cost Function of Neural Networks

# Gradient Computation

## Cost Function:

$$J(\Theta) = -\frac{1}{m} \left( \sum_{i=1}^m \sum_{k=1}^K y_k^{(i)} \log(h_{\Theta}(x^{(i)}))_k + (1 - y_k^{(i)}) \log(1 - h(x^{(i)}; \theta))_k \right) + \frac{\lambda}{2m} \sum_{l=1}^{L-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} (\Theta_{ij}^{(l)})^2$$

## Optimization Objective:

$$\min_{\Theta} J(\Theta)$$

# Gradient Computation

## Cost Function:

$$J(\theta) = -\frac{1}{m} \left( \sum_{i=1}^m \sum_{k=1}^K y_k^{(i)} \log(h_{\theta}(x^{(i)}))_k + (1 - y_k^{(i)}) \log(1 - h(x^{(i)}; \theta))_k \right) + \frac{\lambda}{2m} \sum_{l=1}^{L-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} (\theta_{ij}^{(l)})^2$$

## Optimization Objective:

$$\min_{\theta} J(\theta)$$

## 2-step Strategy:

Compute	Dimension	How?
---------	-----------	------

$J(\theta)$	$ \theta_{ij}^l  = \sum_{j=1}^{L-1} s_{j+1} \times s_j + 1$	<b>Forward Propagation</b>
$\frac{\partial}{\partial \theta_{ij}^{(l)}} J(\theta)$	$\theta_{ij}^l \in \mathbb{R}$	

# Gradient Computation

## Cost Function:

$$J(\theta) = -\frac{1}{m} \left( \sum_{i=1}^m \sum_{k=1}^K y_k^{(i)} \log(h_{\theta}(x^{(i)}))_k + (1 - y_k^{(i)}) \log(1 - h(x^{(i)}; \theta))_k \right) + \frac{\lambda}{2m} \sum_{l=1}^{L-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} (\theta_{ij}^{(l)})^2$$

## Optimization Objective:

$$\min_{\theta} J(\theta)$$

## 2-step Strategy:

Compute	Dimension	How?
$J(\theta)$	$ \theta_{ij}^l  = \sum_{j=1}^{L-1} s_{j+1} \times s_j + 1$	Forward Propagation
$\frac{\partial}{\partial \theta_{ij}^{(l)}} J(\theta)$	$\theta_{ij}^l \in \mathbb{R}$	Back Propagation



# Gradient Computation


To understand how **back propagation** works to compute the partial derivative terms  $\frac{\partial}{\partial \Theta_{ij}^{(l)}} J(\Theta)$ , let's focus on a simple example involving just **one** training example,  $(x, y)$ .

# Gradient Computation

Given  $(x, y)$ :

## 1. Forward Propagation

$$a^{(1)} = x$$



# Gradient Computation


Given  $(x, y)$ :

## 1. Forward Propagation

$$a^{(1)} = x$$

$$z^{(2)} = \Theta^{(1)} a^{(1)}$$

$$a^{(2)} = g(z^{(2)})$$



# Gradient Computation

Given  $(x, y)$ :

## 1. Forward Propagation


$$a^{(1)} = x$$

$$z^{(2)} = \Theta^{(1)} a^{(1)}$$

$$a^{(2)} = g(z^{(2)})$$

$$z^{(3)} = \Theta^{(2)} a^{(2)}$$

$$a^{(3)} = g(z^{(3)})$$



# Gradient Computation

Given  $(x, y)$ :

## 1. Forward Propagation

$$a^{(1)} = x$$

$$z^{(2)} = \Theta^{(1)} a^{(1)}$$


$$a^{(2)} = g(z^{(2)})$$

$$z^{(3)} = \Theta^{(2)} a^{(2)}$$

$$a^{(3)} = g(z^{(3)})$$

$$z^{(4)} = \Theta^{(3)} a^{(3)}$$

$$a^{(4)} = g(z^{(4)}) = h(x^{(i)}; \theta)$$



Forward propagation computes the activation values for all neurons in the network.


# Gradient Computation

Given  $(x, y)$ :

## 2. Back Propagation

**Idea:** Compute the 'error'  $\delta_j^{(l)}$  of node  $j$  in layer  $l$

$$\delta_j^{(4)} = a_j^{(4)} - y_j$$



# Gradient Computation


Given  $(x, y)$ :

## 2. Back Propagation

**Idea:** Compute the 'error'  $\delta_j^{(l)}$  of node  $j$  in layer  $l$

$\delta_j^{(3)}$  = compute errors for layer 3 nodes

$$\delta_j^{(4)} = a_j^{(4)} - y_j$$



# Gradient Computation

Given  $(x, y)$ :


## 2. Back Propagation

**Idea:** Compute the 'error'  $\delta_j^{(l)}$  of node  $j$  in layer  $l$

$\delta_j^{(2)}$  = compute errors for layer 2 nodes

$\delta_j^{(3)}$  = compute errors for layer 3 nodes

$\delta_j^{(4)} = a_j^{(4)} - y_j$





# Gradient Computation

Given  $(x, y)$ :

## 2. Back Propagation


**Idea:** Compute the 'error'  $\delta_j^{(l)}$  of node  $j$  in layer  $l$

Input  $a^{(1)}$  is without error

$\delta_j^{(2)}$  = compute errors for layer 2 nodes

$\delta_j^{(3)}$  = compute errors for layer 3 nodes

$\delta_j^{(4)} = a_j^{(4)} - y_j$



*Backprop computes the error associated with each activation in the network.*


# Gradient Computation

Given  $(x, y)$ :

## 2. Back Propagation

**Example:** Compute the 'error'  $\delta_j^{(4)}$ :

$$\delta_j^{(4)} = a_j^{(4)} - y_j$$




# Gradient Computation

Given  $(x, y)$ :

## 2. Back Propagation

**Example:** Compute the 'error'  $\delta_j^{(4)}$ :

$$\begin{aligned}\delta_j^{(4)} &= a_j^{(4)} - y_j \\ &= h(x^{(i)}; \theta)_j - y_j \\ &= j^{\text{th}} \text{ prediction} - y_j\end{aligned}$$



# Gradient Computation

Given  $(x, y)$ :


## 2. Back Propagation

**Example:** Compute the 'error'  $\delta_j^{(4)}$ :

$$\begin{aligned}\delta_j^{(4)} &= a_j^{(4)} - y_j \\ &= h(x^{(i)}; \theta)_j - y_j \\ &= j^{\text{th}} \text{ prediction} - y_j\end{aligned}$$

**Vectorization:**

$$\delta^{(4)} = a^{(4)} - y$$




# Gradient Computation

Given  $(x, y)$ :

## 2. Back Propagation

**Vectorization:** Compute all 'error'  
 $\delta^{(l)}$ :

$$\delta^{(4)} = a^{(4)} - y$$



# Gradient Computation


Given  $(x, y)$ :

## 2. Back Propagation

**Vectorization:** Compute all 'error'  
 $\delta^{(l)}$ :

$$\delta^{(4)} = a^{(4)} - y$$

$$\delta^{(3)} = (\Theta^{(3)})^T \delta^{(4)} .* [g'(z^{(3)})]$$



# Gradient Computation

Given  $(x, y)$ :


## 2. Back Propagation

**Vectorization:** Compute all 'error'  
 $\delta^{(l)}$ :

$$\delta^{(4)} = a^{(4)} - y$$

$$\delta^{(3)} = (\Theta^{(3)})^T \delta^{(4)} .* [g'(z^{(3)})]$$

$$\delta^{(2)} = (\Theta^{(2)})^T \delta^{(3)} .* [g'(z^{(2)})]$$



# Gradient Computation

Given  $(x, y)$ :


## 2. Back Propagation

**Vectorization:** Compute all 'error'  
 $\delta^{(l)}$ :

$$\delta^{(4)} = a^{(4)} - y$$

$$\delta^{(3)} = (\theta^{(3)})^T \delta^{(4)} .* [g'(z^{(3)})]$$

$$\delta^{(2)} = (\theta^{(2)})^T \delta^{(3)} .* [g'(z^{(2)})]$$



$g'(z)$  is the derivative of the activation function  $g(\cdot)$  evaluated at the input values given by  $z$ .

$.*$  refers to dot-product multiplication.



# Gradient Computation

Given  $(x, y)$ :


## 2. Back Propagation

**Vectorization:** Compute all 'error'  $\delta^{(l)}$ :

$$\delta^{(4)} = a^{(4)} - y$$

$$\delta^{(3)} = (\Theta^{(3)})^T \delta^{(4)} .* [a^{(3)} .* (1 - a^{(3)})]$$

$$\delta^{(2)} = (\Theta^{(2)})^T \delta^{(3)} .* [a^{(2)} .* (1 - a^{(2)})]$$



# Gradient Computation

Given  $(x, y)$ :


## 2. Back Propagation

**Vectorization:** Compute all 'error'  $\delta^{(l)}$ :

$$\delta^{(4)} = a^{(4)} - y$$

$$\delta^{(3)} = (\Theta^{(3)})^T \delta^{(4)} .* [a^{(3)} .* (1 - a^{(3)})]$$

$$\delta^{(2)} = (\Theta^{(2)})^T \delta^{(3)} .* [a^{(2)} .* (1 - a^{(2)})]$$



The proof is complicated, but it can be shown that

$$(\Theta^{(l)})^T \delta^{(l+1)} .* [a^{(l)} .* (1 - a^{(l)})]$$

is **equivalent** to simply calculating

$$a^{(l)} \delta^{(l+1)}$$

# Gradient Computation


Given  $(x, y)$ :

## 2. Back Propagation

Upshot:

$$\frac{\partial}{\partial \Theta_{ij}^{(l)}} J(\Theta) = a_j^{(l)} \delta_i^{(l+1)}$$

(if we ignore  $\lambda$  or simply set  $\lambda = 0$ )



So much for back propagation for one training example  $(x, y)$ .

So much for back propagation for one training example  $(x, y)$ .

What about lots of training examples?

# Gradient Computation

Given  $\{(x^{(1)}, y^{(1)}), \dots, (x^{(m)}, y^{(m)})\}$ :

## 1. Forward Propagation

$$a^{(1)} = x$$

$$z^{(2)} = \Theta^{(1)} a^{(1)}$$


$$a^{(2)} = g(z^{(2)})$$

$$z^{(3)} = \Theta^{(2)} a^{(2)}$$

$$a^{(3)} = g(z^{(3)})$$

$$z^{(4)} = \Theta^{(3)} a^{(3)}$$

$$a^{(4)} = g(z^{(4)}) = h(x^{(i)}; \theta)$$



# Back Propagation Algorithm

## Algorithm: Backpropagation

```
» Given  $\{(x^{(1)}, y^{(1)}), \dots, (x^{(m)}, y^{(m)})\}$ 
» Set  $\Delta_{i,j}^l = 0$ 
  For  $i = 1 : m$ 
    Set  $a^{(1)} = x^{(i)}$ 
    Perform forward propagation to compute  $a^{(l)}$  for  $l = 2, 3, \dots, L$ 
    Using  $y^{(i)}$ , compute  $\delta^{(L)} = a^{(L)} - y^{(i)}$ 
    Compute  $\delta^{(L-1)}, \delta^{(L-2)}, \dots, \delta^{(2)}$ 
     $\Delta_{i,j}^{(l)} := \Delta_{ij}^{(l)} + a_j^{(l)} \delta_i^{(l+1)}$ 
  end
```

# Back Propagation Algorithm

## Algorithm: Backpropagation

```
» Given  $\{(x^{(1)}, y^{(1)}), \dots, (x^{(m)}, y^{(m)})\}$ 
» Set  $\Delta_{i,j}^l = 0$ 
  For  $i = 1 : m$ 
    Set  $a^{(1)} = x^{(i)}$ 
    Perform forward propagation to compute  $a^{(l)}$  for  $l = 2, 3, \dots, L$ 
    Using  $y^{(i)}$ , compute  $\delta^{(L)} = a^{(L)} - y^{(i)}$ 
    Compute  $\delta^{(L-1)}, \delta^{(L-2)}, \dots, \delta^{(2)}$ 
     $\Delta_{i,j}^{(l)} := \Delta_{ij}^{(l)} + a_j^{(l)} \delta_i^{(l+1)}$  # Vectorized:  $\Delta^{(l)} := \Delta^{(l)} + \delta^{(l+1)} (a^{(l)})^T$ 
  end
```



# Back Propagation Algorithm

## Algorithm: Backpropagation


- » Given  $\{(x^{(1)}, y^{(1)}), \dots, (x^{(m)}, y^{(m)})\}$
- » Set  $\Delta_{i,j}^l = 0$ 
  - For  $i = 1 : m$ 
    - Set  $a^{(1)} = x^{(i)}$
    - Perform forward propagation to compute  $a^{(l)}$  for  $l = 2, 3, \dots, L$
    - Using  $y^{(i)}$ , compute  $\delta^{(L)} = a^{(L)} - y^{(i)}$
    - Compute  $\delta^{(L-1)}, \delta^{(L-2)}, \dots, \delta^{(2)}$
    - $\Delta_{i,j}^{(l)} := \Delta_{ij}^{(l)} + a_j^{(l)} \delta_i^{(l+1)}$  # Vectorized:  $\Delta^{(l)} := \Delta^{(l)} + \delta^{(l+1)} (a^{(l)})^T$
  - end
- » Compute partial derivatives  $\frac{\partial}{\partial \Theta_{ij}^{(l)}} J(\Theta) = D_{ij}^{(l)}$ , by

$$D_{ij}^{(l)} := \frac{1}{m} \Delta_{ij}^{(l)} + \lambda \Theta_{ij}^{(l)} \quad \text{if } j \neq 0$$


$$D_{ij}^{(l)} := \frac{1}{m} \Delta_{ij}^{(l)} \quad \text{if } j = 0$$

# Backpropagation Intuition


## Forward Propagation




## Forward Propagation




## Forward Propagation




## Forward Propagation



## Forward Propagation




## Forward Propagation







## Forward Propagation




## Forward Propagation




## Forward Propagation




## Forward Propagation




## Forward Propagation




## Forward Propagation



## Forward Propagation




## Forward Propagation







## Forward Propagation



## Forward Propagation



## Forward Propagation



Forward propagation flows from *Left-to-Right*

Back propagation flows from *Right-to-Left*

# What back propagation is doing

**Neural Network Cost Function for  $h(x^{(i)}; \Theta) \in \mathbb{R}^K$**

$$J(\Theta) = -\frac{1}{m} \left( \sum_{i=1}^m \sum_{k=1}^K y_k^{(i)} \log(h(x^{(i)}; \Theta)_k) + (1 - y_k^{(i)}) \log(1 - h(x^{(i)}; \Theta)_k) \right) + \frac{\lambda}{2m} \left( \sum_{l=1}^{L-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} \Theta_{ij}^{(l)} \right)^2$$

# What back propagation is doing

**Neural Network Cost Function for**  $h(x^{(i)}; \Theta) \in \mathbb{R}^K$

$$J(\Theta) = -\frac{1}{m} \left( \sum_{i=1}^m \sum_{k=1}^K y_k^{(i)} \log(h(x^{(i)}; \Theta)_k) + (1 - y_k^{(i)}) \log(1 - h(x^{(i)}; \Theta)_k) \right) + \frac{\lambda}{2m} \left( \sum_{l=1}^{L-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} \Theta_{ij}^{(l)} \right)^2$$

*...without regularization:*

$$J(\Theta) = -\frac{1}{m} \left( \sum_{i=1}^m \sum_{k=1}^K y_k^{(i)} \log(h(x^{(i)}; \Theta)_k) + (1 - y_k^{(i)}) \log(1 - h(x^{(i)}; \Theta)_k) \right)$$

# What back propagation is doing

**Neural Network Cost Function for  $h(x^{(i)}; \Theta) \in \mathbb{R}^K$**

$$J(\Theta) = -\frac{1}{m} \left( \sum_{i=1}^m \sum_{k=1}^K y_k^{(i)} \log(h(x^{(i)}; \Theta)_k) + (1 - y_k^{(i)}) \log(1 - h(x^{(i)}; \Theta)_k) \right) + \frac{\lambda}{2m} \left( \sum_{l=1}^{L-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} \Theta_{ij}^{(l)} \right)^2$$

*...without regularization:*

$$J(\Theta) = -\frac{1}{m} \left( \sum_{i=1}^m \sum_{k=1}^K y_k^{(i)} \log(h(x^{(i)}; \Theta)_k) + (1 - y_k^{(i)}) \log(1 - h(x^{(i)}; \Theta)_k) \right)$$

*...for only one output unit ( $K = 1$ ):*

$$J(\Theta) = -\frac{1}{m} \left( \sum_{i=1}^m y^{(i)} \log(h(x^{(i)}; \Theta)) + (1 - y^{(i)}) \log(1 - h(x^{(i)}; \Theta)) \right)$$

# What back propagation is doing

**Neural Network Cost Function for**  $h(x^{(i)}; \Theta) \in \mathbb{R}^K$

$$J(\Theta) = -\frac{1}{m} \left( \sum_{i=1}^m \sum_{k=1}^K y_k^{(i)} \log(h(x^{(i)}; \Theta)_k) + (1 - y_k^{(i)}) \log(1 - h(x^{(i)}; \Theta)_k) \right) + \frac{\lambda}{2m} \left( \sum_{l=1}^{L-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} \Theta_{ij}^{(l)} \right)^2$$

*...without regularization:*

$$J(\Theta) = -\frac{1}{m} \left( \sum_{i=1}^m \sum_{k=1}^K y_k^{(i)} \log(h(x^{(i)}; \Theta)_k) + (1 - y_k^{(i)}) \log(1 - h(x^{(i)}; \Theta)_k) \right)$$

*...for only one output unit ( $K = 1$ ):*

$$J(\Theta) = -\frac{1}{m} \left( \sum_{i=1}^m y^{(i)} \log(h(x^{(i)}; \Theta)) + (1 - y^{(i)}) \log(1 - h(x^{(i)}; \Theta)) \right)$$

*...for one particular training example:  $(x^{(i)}, y^{(i)})$*



# What back propagation is doing

*...for one particular training example:  $(x^{(i)}, y^{(i)})$*

$$J(\Theta) = -\frac{1}{m} \left( \sum_{i=1}^m y^{(i)} \log(h(x^{(i)}; \Theta)) + (1 - y^{(i)}) \log(1 - h(x^{(i)}; \Theta)) \right)$$

# What back propagation is doing

*...for one particular training example:  $(x^{(i)}, y^{(i)})$*

$$J(\Theta) = -\frac{1}{m} \left( \sum_{i=1}^m y^{(i)} \log(h(x^{(i)}; \Theta)) + (1 - y^{(i)}) \log(1 - h(x^{(i)}; \Theta)) \right)$$

*...which calculates the cost associated with the  $i$ th training example:*

$$\text{cost}(i) = y^{(i)} \log(h(x^{(i)}; \Theta)) + (1 - y^{(i)}) \log(1 - h(x^{(i)}; \Theta))$$

# What back propagation is doing

*...for one particular training example:  $(x^{(i)}, y^{(i)})$*


$$J(\Theta) = -\frac{1}{m} \left( \sum_{i=1}^m y^{(i)} \log(h(x^{(i)}; \Theta)) + (1 - y^{(i)}) \log(1 - h(x^{(i)}; \Theta)) \right)$$

*...which calculates the cost associated with the  $i$ th training example:*

$$\text{cost}(i) = y^{(i)} \log(h(x^{(i)}; \Theta)) + (1 - y^{(i)}) \log(1 - h(x^{(i)}; \Theta))$$

**Question:** *How well is the network doing on example  $i$ ?*


## Backward Propagation



$\delta_j^{(l)}$  = informally, the error of **cost** for activation  $a_j^{(l)}$  of node  $j$  in layer  $l$ .

formally,  $\delta_j^{(l)}$  is the partial derivative with respect to input  $z_j^{(l)}$  of **cost** of training example  $i$ :

## Backward Propagation




$\delta_j^{(l)}$  = informally, the error of **cost** for activation  $a_j^{(l)}$  of node  $j$  in layer  $l$ .

formally,  $\delta_j^{(l)}$  is the partial derivative with respect to input  $z_j^{(l)}$  of **cost** of training example  $i$ :

$$\delta_j^{(l)} = \frac{\partial}{\partial z_j^{(l)}} \text{cost}(i), \quad \text{where } \text{cost}(i) = y^{(i)} \log h(x^{(i)}; \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}; \theta))$$

## Backward Propagation




$\delta_j^{(l)}$  = informally, the error of **cost** for activation  $a_j^{(l)}$  of node  $j$  in layer  $l$ .

formally,  $\delta_j^{(l)}$  is the partial derivative with respect to input  $z_j^{(l)}$  of **cost** of training example  $i$ :

$$\delta_j^{(l)} = \frac{\partial}{\partial z_j^{(l)}} \text{cost}(i), \quad \text{where } \text{cost}(i) = y^{(i)} \log h(x^{(i)}; \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}; \theta))$$

Intuitively, if we go into the network and change these  $z_j^{(l)}$  values, that will change the **cost**.

## Backward Propagation




$\delta_j^{(l)}$  = the 'error' for  $a_j^{(l)}$

$$\delta_1^{(4)} = a_1^{(4)} - y^{(i)}$$

# Difference between the actual value of  $y^{(i)}$  and the predicted value of  $y^{(i)}$

## Backward Propagation




$\delta_j^{(l)}$  = the 'error' for  $a_j^{(l)}$

$\delta_1^{(4)} = a_1^{(4)} - y^{(i)}$  # Difference between the actual value of  $y^{(i)}$  and the predicted value of  $y^{(i)}$



## Backward Propagation




$\delta_j^{(l)}$  = the 'error' for  $a_j^{(l)}$

$$\delta_1^{(4)} = a_1^{(4)} - y^{(i)}$$


# Difference between the actual value of  $y^{(i)}$  and the predicted value of  $y^{(i)}$

## Backward Propagation




Let's look at how to calculate  $\delta_2^{(2)}$ .

## Backward Propagation




Let's look at how to calculate  $\delta_2^{(2)}$ .

## Back Propagation



## Back Propagation



## How to Compute $\delta_2^{(2)}$ :


$$(\Theta^{(l)})^T \delta^{(l+1)} .* [a^{(l)} .* (1 - a^{(l)})]$$

$$\delta_2^{(2)} = \Theta_{12}^{(2)} \delta_1^{(3)} .* (1 - a_1^{(2)}) + \Theta_{22}^{(2)} \delta_2^{(3)} .* (1 - a_2^{(2)})]$$

$$\Theta^{(2)} = \begin{bmatrix} \Theta_{10}^{(2)} & \Theta_{11}^{(2)} & \Theta_{12}^{(2)} \\ \Theta_{20}^{(2)} & \Theta_{21}^{(2)} & \Theta_{22}^{(2)} \end{bmatrix}$$

where  $\Theta^{(2)} \in \mathbb{R}^{2 \times 3}$

## Back Propagation



## How to Compute $\delta_2^{(2)}$ :

$$(\Theta^{(l)})^T \delta^{(l+1)} .* [a^{(l)} .* (1 - a^{(l)})]$$

$$\delta_2^{(2)} = \theta_{12}^{(2)} \delta_1^{(3)} .* (1 - a_1^{(2)}) + \theta_{22}^{(2)} \delta_2^{(3)} .* (1 - a_2^{(2)})]$$

$$\Theta^{(2)} = \begin{bmatrix} \theta_{10}^{(2)} & \theta_{11}^{(2)} & \theta_{12}^{(2)} \\ \theta_{20}^{(2)} & \theta_{21}^{(2)} & \theta_{22}^{(2)} \end{bmatrix}$$

where  $\Theta^{(2)} \in \mathbb{R}^{2 \times 3}$

Informally,  $\delta_2^{(2)}$  is the weighted sum of the errors  $\delta_1^{(3)}$  and  $\delta_2^{(3)}$ , where the weights  $\theta_{12}^{(2)}$  and  $\theta_{22}^{(2)}$  are the corresponding edge strengths.



# Implementation Details

1. Unrolling (or unstacking) Matrices
2. Random Initialization of Parameter Matrices



# Unroll and Reshape

## Parameter Matrices Dimensions

$$\Theta^{(1)} \in \mathbb{R}^{10 \times 11} \quad \Theta^{(2)} \in \mathbb{R}^{10 \times 11} \quad \Theta^{(3)} \in \mathbb{R}^{1 \times 11}$$

## Derivative Matrices Dimensions

$$D^{(1)} \in \mathbb{R}^{10 \times 11} \quad D^{(2)} \in \mathbb{R}^{10 \times 11} \quad D^{(3)} \in \mathbb{R}^{1 \times 11}$$

# Unroll and Reshape

## Parameter Matrices Dimensions

$$\Theta^{(1)} \in \mathbb{R}^{10 \times 11} \quad \Theta^{(2)} \in \mathbb{R}^{10 \times 11} \quad \Theta^{(3)} \in \mathbb{R}^{1 \times 11}$$

## Derivative Matrices Dimensions

$$D^{(1)} \in \mathbb{R}^{10 \times 11} \quad D^{(2)} \in \mathbb{R}^{10 \times 11} \quad D^{(3)} \in \mathbb{R}^{1 \times 11}$$

## Python

```
# Format training set
X_train = X_train.reshape(60000, 784)
X_train = X_train.astype('float32')
X_train = X_train/255

# Format test set
X_test = X_test.reshape(10000, 784)
X_test = X_test.astype('float32')
X_test = X_test/255

print("Training matrix shape", X_train.shape)
print("Testing matrix shape", X_test.shape)
```

Each of the 60,000 training examples in the MNIST dataset is a 28x28 matrix.

`X_*.reshape(60000, 784)` takes 60,000 examples of `X_*` and reshapes into a 784-dimension vector.

## Example

### Network Dimensions

$$s_1 = s_2 = 10; s_3 = 1$$

### Parameter Matrices Dimensions

$$\Theta^{(1)} \in \mathbb{R}^{10 \times 11}$$

$$\Theta^{(2)} \in \mathbb{R}^{10 \times 11}$$


$$\Theta^{(3)} \in \mathbb{R}^{1 \times 11}$$

### Derivative Matrices Dimensions

$$D^{(1)} \in \mathbb{R}^{10 \times 11}$$

$$D^{(2)} \in \mathbb{R}^{10 \times 11}$$

$$D^{(3)} \in \mathbb{R}^{1 \times 11}$$



Random Initialization

# Initializing Theta

## Initial value of $\Theta$ :

To implement either gradient descent or advanced optimization methods, such as `bfgs`, you need to specify **initial values for  $\Theta$** .

# Initializing Theta

## Initial value of $\Theta$ :

To implement either gradient descent or advanced optimization methods, such as **bfgs**, you need to specify **initial values for  $\Theta$** .

Recall that for logistic regression, we simply initialized theta to a **vector of zeros**.

**Why?** Because  $g(z) = 1$  if  $z \geq 0$  and  $g(z) = 0$  if  $z < 0$

# Initializing Theta

## Initial value of $\Theta$ :

To implement either gradient descent or advanced optimization methods, such as `bfgs`, you need to specify **initial values for  $\Theta$** .


Recall that for logistic regression, we simply initialized theta to a **vector of zeros**.

**Why?** Because  $g(z) = 1$  if  $z \geq 0$  and  $g(z) = 0$  if  $z < 0$

`initialTheta = zeros(n,1)` does **not** work for neural networks.

**Why not?**


# Zero initialization of $\Theta$



Suppose  $\theta_{ij}^{(l)} = 0$ , for all  $i, j, l$ .



# Zero initialization of $\Theta$




Suppose  $\Theta_{ij}^{(l)} = 0$ , for all  $i, j, l$ .

So:

$$\Theta_{10}^{(1)} = \Theta_{20}^{(1)}$$

# Zero initialization of $\Theta$




Suppose  $\Theta_{ij}^{(l)} = 0$ , for all  $i, j, l$ .

So:

$$\Theta_{10}^{(1)} = \Theta_{20}^{(1)}$$

$$\Theta_{11}^{(1)} = \Theta_{21}^{(1)}$$

# Zero initialization of $\Theta$



Suppose  $\theta_{ij}^{(l)} = 0$ , for all  $i, j, l$ .


So:

$$\theta_{10}^{(1)} = \theta_{20}^{(1)}$$

$$\theta_{11}^{(1)} = \theta_{21}^{(1)}$$

$$\theta_{12}^{(1)} = \theta_{22}^{(1)}$$

# Zero initialization of $\Theta$



Suppose  $\theta_{ij}^{(l)} = 0$ , for all  $i, j, l$ .

So:


Hence,

$$\theta_{10}^{(1)} = \theta_{20}^{(1)} \quad a_1^{(2)} = a_2^{(2)}$$

$$\theta_{11}^{(1)} = \theta_{21}^{(1)}$$

$$\theta_{12}^{(1)} = \theta_{22}^{(1)}$$

# Zero initialization of $\Theta$



Suppose  $\Theta_{ij}^{(l)} = 0$ , for all  $i, j, l$ .

So:


Hence,

$$\Theta_{10}^{(1)} = \Theta_{20}^{(1)} \quad a_1^{(2)} = a_2^{(2)}$$

$$\Theta_{11}^{(1)} = \Theta_{21}^{(1)} \quad \text{and}$$

$$\Theta_{12}^{(1)} = \Theta_{22}^{(1)} \quad \delta_1^{(2)} = \delta_2^{(2)}$$

# Zero initialization of $\Theta$



Suppose  $\Theta_{ij}^{(l)} = 0$ , for all  $i, j, l$ .

So:

Hence,

$$\Theta_{10}^{(1)} = \Theta_{20}^{(1)} \quad a_1^{(2)} = a_2^{(2)}$$


$$\Theta_{11}^{(1)} = \Theta_{21}^{(1)} \quad \text{and}$$

$$\Theta_{12}^{(1)} = \Theta_{22}^{(1)} \quad \delta_1^{(2)} = \delta_2^{(2)}$$

Furthermore,

$$\frac{\partial}{\partial \Theta_{10}^{(1)}} J(\Theta) = \frac{\partial}{\partial \Theta_{20}^{(1)}} J(\Theta), \text{ so } \Theta_{10}^{(1)} = \Theta_{20}^{(1)}$$

# Zero initialization of $\Theta$



**Upshot:** Since  $\delta_1^{(l)} = \delta_2^{(l)} = 0$ , after each update, the parameters associated with each input that go into each of the two hidden units are identical.

Suppose  $\Theta_{ij}^{(l)} = 0$ , for all  $i, j, l$ .

So:

Hence,

$$\Theta_{10}^{(1)} = \Theta_{20}^{(1)} \quad a_1^{(2)} = a_2^{(2)}$$

$$\Theta_{11}^{(1)} = \Theta_{21}^{(1)} \quad \text{and}$$


$$\Theta_{12}^{(1)} = \Theta_{22}^{(1)} \quad \delta_1^{(2)} = \delta_2^{(2)}$$

Furthermore,

$$\frac{\partial}{\partial \Theta_{10}^{(1)}} J(\Theta) = \frac{\partial}{\partial \Theta_{20}^{(1)}} J(\Theta), \text{ so } \Theta_{10}^{(1)} = \Theta_{20}^{(1)}$$

$$\frac{\partial}{\partial \Theta_{11}^{(1)}} J(\Theta) = \frac{\partial}{\partial \Theta_{21}^{(1)}} J(\Theta), \text{ so } \Theta_{11}^{(1)} = \Theta_{21}^{(1)}$$

# Zero initialization of $\Theta$



**Upshot:** Since  $\delta_1^{(l)} = \delta_2^{(l)} = 0$ , after each update, the parameters associated with each input that go into each of the two hidden units are identical.

So, the two hidden units are *still* computing the same function as the input:

$$a_1^{(2)} = a_2^{(2)}$$

Suppose  $\Theta_{ij}^{(l)} = 0$ , for all  $i, j, l$ .

So:

Hence,

$$\Theta_{10}^{(1)} = \Theta_{20}^{(1)} \quad a_1^{(2)} = a_2^{(2)}$$

$$\Theta_{11}^{(1)} = \Theta_{21}^{(1)} \quad \text{and}$$

$$\Theta_{12}^{(1)} = \Theta_{22}^{(1)} \quad \delta_1^{(2)} = \delta_2^{(2)}$$

Furthermore,

$$\frac{\partial}{\partial \Theta_{10}^{(1)}} J(\Theta) = \frac{\partial}{\partial \Theta_{20}^{(1)}} J(\Theta), \text{ so } \Theta_{10}^{(1)} = \Theta_{20}^{(1)}$$

$$\frac{\partial}{\partial \Theta_{11}^{(1)}} J(\Theta) = \frac{\partial}{\partial \Theta_{21}^{(1)}} J(\Theta), \text{ so } \Theta_{11}^{(1)} = \Theta_{21}^{(1)}$$

$$\frac{\partial}{\partial \Theta_{12}^{(1)}} J(\Theta) = \frac{\partial}{\partial \Theta_{22}^{(1)}} J(\Theta), \text{ so } \Theta_{12}^{(1)} = \Theta_{22}^{(1)}$$



# Random Initialization of $\Theta$

To get around the previous problem with **zero-initialization**, **random initialization** of  $\Theta$  is a technique that assigns each value  $\Theta_{ij}^{(l)}$  a random scalar value in some range,  $[-\varepsilon, \varepsilon]$ .<sup>1</sup>

---

<sup>1</sup>Note that this use of the variable  $\varepsilon$  is entirely different than its use in numerical gradient checking.

# Random Initialization of $\Theta$




To get around the **problem of symmetric weights**, initialize each  $\Theta_{ij}^{(l)}$  to a random small values in  $[-\varepsilon, \varepsilon]$  close to zero, that is:

$$-\varepsilon \leq \Theta_{ij}^{(l)} \leq \varepsilon$$






# Training a Neural Network

Pick a network architecture (i.e., the connectivity pattern between neurons)






# Training a Neural Network

Pick a network architecture (i.e., the connectivity pattern between neurons)






# Training a Neural Network

Pick a network architecture (i.e., the connectivity pattern between neurons)






# Training a Neural Network

Pick a network architecture (i.e., the connectivity pattern between neurons)



# Training a Neural Network

Pick a network architecture (i.e., the connectivity pattern between neurons)






**Number of input units:**



# Training a Neural Network

Pick a network architecture (i.e., the connectivity pattern between neurons)






**Number of input units:**

*dimension of features:*  $x^{(i)}$

# Training a Neural Network

Pick a network architecture (i.e., the connectivity pattern between neurons)






**Number of input units:**

*dimension of features:*  $x^{(i)}$

**Number of output units:**

# Training a Neural Network

Pick a network architecture (i.e., the connectivity pattern between neurons)



**Number of input units:**

*dimension of features:*

$x^{(i)}$




**Number of output units:**

*number of classes:*

$K$  is a  $\mathbb{R}^K$  vector

# Training a Neural Network

Pick a network architecture (i.e., the connectivity pattern between neurons)



**Number of input units:**

**Number of output units:**

**Num. of hidden layers:**

*dimension of features:*




*number of classes:*

$x^{(i)}$

$K$  is a  $\mathbb{R}^K$  vector

# Training a Neural Network

Pick a network architecture (i.e., the connectivity pattern between neurons)



**Number of input units:**

**Number of output units:**

**Num. of hidden layers:**

*dimension of features:*

*number of classes:*

*default to try:*




$x^{(i)}$

$K$  is a  $\mathbb{R}^K$  vector

1

# Training a Neural Network

Pick a network architecture (i.e., the connectivity pattern between neurons)



**Number of input units:**

**Number of output units:**

**Num. of hidden layers:**

**Num. hidden units:**

*dimension of features:*

*number of classes:*

*default to try:*

*if  $> 1$  hidden layers*

$x^{(i)}$




$K$  is a  $\mathbb{R}^K$  vector

1

*all layers same Number*

# Training a Neural Network

Pick a network architecture (i.e., the connectivity pattern between neurons)



**Number of input units:**

**Number of output units:**

**Num. of hidden layers:**

**Num. hidden units:**

*dimension of features:*

*number of classes:*

*default to try:*

*if  $> 1$  hidden layers  
more hidden units better*

$x^{(i)}$

$K$  is a  $\mathbb{R}^K$  vector

1

*all layers same Number*

*but more units,  
more computation*

# Training a Neural Network

0. Pick a network architecture
1. Randomly **initialize weights**
2. Implement **forward propagation** to compute  $h(x^{(i)}; \theta)$  for any  $x^{(i)}$ .
3. Implement code to compute the **cost function**  $J(\theta)$
4. Implement **back propagation** to compute partial derivatives  $\frac{\partial}{\partial \theta_{jk}^{(l)}} J(\theta)$



# Training a Neural Network

0. Pick a network architecture


1. Randomly **initialize weights**

2. Implement **forward propagation** to compute  $h(x^{(i)}; \theta)$  for any  $x^{(i)}$ .

3. Implement code to compute the **cost function**  $J(\theta)$

4. Implement **back propagation** to compute partial derivatives  $\frac{\partial}{\partial \theta_{jk}^{(l)}} J(\theta)$

```
for i = 1:m {  
  for each training example  $(x^{(i)}, y^{(i)})$ :  
    execute forward prop to get activations  $a^{(l)}$   
    and back prop to get delta terms  $\delta^{(l)}$   
    (for  $l = 2, 3, \dots, L$ )  
    compute  $\Delta^{(l)} := \Delta^{(l)} + \delta^{(l+1)} \cdot (a^{(l)})^T$   
  }  
  compute  $\frac{\partial}{\partial \theta_{jk}^{(l)}} J(\theta)$ 
```



# Training a Neural Network

0. Pick a network architecture


1. Randomly **initialize weights**

2. Implement **forward propagation** to compute  $h(x^{(i)}; \theta)$  for any  $x^{(i)}$ .

3. Implement code to compute the **cost function**  $J(\theta)$

4. Implement **back propagation** to compute partial derivatives  $\frac{\partial}{\partial \theta_{jk}^{(l)}} J(\theta)$

```
for i = 1:m {  
    for each training example  $(x^{(i)}, y^{(i)})$ :  
        execute forward prop to get activations  $a^{(l)}$   
        and back prop to get delta terms  $\delta^{(l)}$   
        (for  $l = 2, 3, \dots, L$ )  
        compute  $\Delta^{(l)} := \Delta^{(l)} + \delta^{(l+1)} \cdot (a^{(l)})^T$   
    }  
    compute  $\frac{\partial}{\partial \theta_{jk}^{(l)}} J(\theta)$ 
```



# Training a Neural Network


5. Use gradient checking to compare  $\frac{\partial}{\partial \theta_{jk}^{(l)}} J(\Theta)$  computed by back propagation and by numerical estimation of the gradient of  $J(\Theta)$ .

Then **disable** gradient checking.

6. Use gradient descent or an advanced optimization method with back propagation to attempt to minimize the cost function  $J(\Theta)$  as a function of parameters  $\Theta$ .

# Cost Function

The cost function  $J(\theta)$  is **non-convex**.



# Cost Function

The cost function  $J(\Theta)$  is **non-convex**.

**Question:** *Why does gradient-descent work **at all** in neural networks despite non-convexity?*

# Cost Function


The cost function  $J(\Theta)$  is **non-convex**.

**Question:** *Why does gradient-descent work **at all** in neural networks despite non-convexity?*

A partial answer is **over-provisioning**: since there are usually many hidden units, there are many different ways that a neural network can approximate the desired input-output relationship and you only need to find one (Carmon and Duchi 2016).


## Coda: Dropout Regularization

# Dropout Intuition





# Dropout Intuition



# Dropout

Dropout temporarily removes nodes from a network.

Actually, what is temporarily "removed" are the links going in and out of randomly selected nodes.

Dropout simulates **sparse activation** from a given layer, which in effect reduces the model capacity of the network.

## Notes:

- Dropout is used to address overfitting.
- The probability is the probability of a unit in a given layer dropping out.
- Nodes are drawn randomly on each pass.
- So, in practice, different nodes will be dropped in different passes.
- Dropout not used at test time. Otherwise, the predictions would be random. (Keras does this automatically).
- Dropout effectively spreads out weights; ensures that the network does not rely on any single feature.
- By spreading out weights, this effectively reduces the squared norm,  $\|\cdot\|_F^2$ .
- Probability for keeping units can be varied by layer.
- **Downside:** the cost function  $J(\cdot)$  is no longer well-defined.
- To plot  $J(\cdot)$ , to check that it is monotonically decreasing, turn off dropout.

## References

Carmon, Y. and J. Duchi (2016).

Gradient descent efficiently finds the cubic-regularized non-convex Newton step.  
In *Workshop on Nonconvex Optimization for Machine Learning (NIPS 2016)*.

Nielsen, M. (2019).

*Neural Networks and Deep Learning*.  
Determination Press.